



(a) Fixed model and inputs



(c) Query and stop



(b) Rank calibration



With probability $\geq 1 - \delta$ over label-query order:
 $C(x; q_{\text{ref}}) \subseteq C(x; U)$ for every x ; mean extra labels $\leq \tau$ on the monitoring batch